

Lab 8 - PCC177/BCC406

REDES NEURAIS E APRENDIZAGEM EM PROFUNDIDADE

Modelos Generativos

Prof. Eduardo e Prof. Pedro

Objetivos:

- Parte I : Compressão com AE
- Parte II : Detecção de anomalias
- Parte III: Redes Generativas Adversariais

Data da entrega : 12/12/23

- Complete o código (marcado com `ToDo`) e quando requisitado, escreva textos diretamente nos notebooks. Onde tiver *None*, substitua pelo seu código.
- Execute todo notebook e salve tudo em um PDF **nomeado** como "NomeSobrenome-Lab.pdf"
- Envie o PDF via google [FORM](#)

Este notebook é baseado em tensorflow e Keras.

✓ Parte I: Autoencoder para redução de dimensionalidade (30pt)

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import tensorflow as tf

from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
from tensorflow.keras import layers, losses
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Model
```

Carrega dataset Fashion MNIST dataset. Cada imagem tem resolução 28x28 pixels.

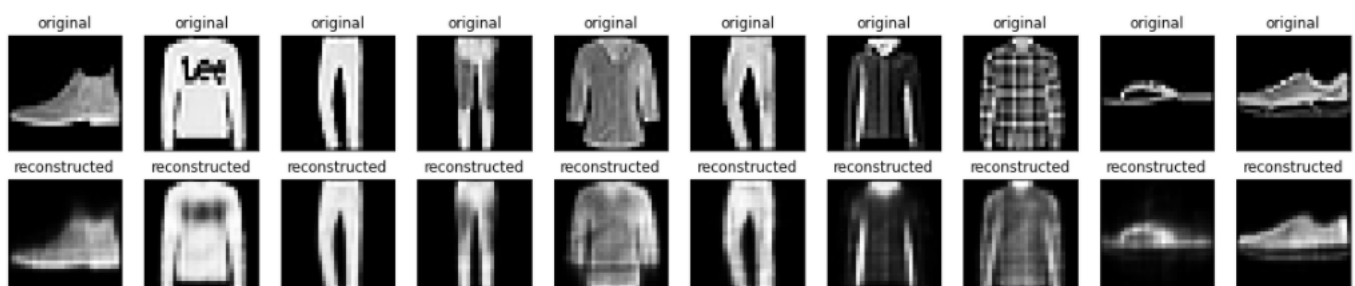
```
(x_train, _), (x_test, _) = fashion_mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

print (x_train.shape)
print (x_test.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz
29515/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz
26421880/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz
5148/5148 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz
4422102/4422102 [=====] - 0s 0us/step
(60000, 28, 28)
(10000, 28, 28)
```

✓ Exemplo de classes



Abaixo exemplo de implementação de autoencoder apenas com camadas densas. O `encoder`, comprime as imagens em 4 dimensões (`latent_dim`), e o `decoder` reconstrói a imagem a partir do vetor latente.

O exemplo abaixo usa a [Keras Model Subclassing API](#).

```
latent_dims = [2, 4, 8, 16, 64]

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
        ])
        self.decoder = tf.keras.Sequential([
            layers.Dense(784, activation='sigmoid'),
            layers.Reshape((28, 28))
        ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded

autoencoders = []

for dim in latent_dims:
    autoencoder = Autoencoder(latent_dim=dim)
    autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())
    autoencoder.fit(x_train, x_train, epochs=10, shuffle=True, validation_data=(x_test, x_test))
    autoencoders.append(autoencoder)
```

Epoch 2/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0254 - val_loss: 0.0242
Epoch 3/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0239 - val_loss: 0.0237
Epoch 4/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0235 - val_loss: 0.0234
Epoch 5/10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0233 - val_loss: 0.0233

```
epoch // 10
1875/1875 [=====] - 6s 3ms/step - loss: 0.0089 - val_loss: 0.0089
Epoch 8/10
1875/1875 [=====] - 8s 4ms/step - loss: 0.0088 - val_loss: 0.0088
Epoch 9/10
1875/1875 [=====] - 5s 3ms/step - loss: 0.0087 - val_loss: 0.0088
Epoch 10/10
1875/1875 [=====] - 7s 4ms/step - loss: 0.0087 - val_loss: 0.0088
```

Treine o modelo e veja os resultados da re-construção.

```
# Exibição das imagens para todos os modelos em uma única figura
n_models = len(latent_dims)
n_images = 10
vertical_space = 2

plt.figure(figsize=(20, (4 + vertical_space) * n_models))

for i, (dim, autoencoder) in enumerate(zip(latent_dims, autoencoders), 1):
    encoded_imgs = autoencoder.encoder(x_test).numpy()
    decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

    #plt.text(n_images // 2, -1, f'Latent Dim: {dim}', ha='center', va='center', fontsize=12)

    for j in range(n_images):
        # display original
        ax = plt.subplot(n_models * 2, n_images, (i-1) * (n_images * 2) + j + 1)
        plt.imshow(x_test[j])
        plt.title("Original")
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

    for j in range(n_images):
        # display reconstruction
        ax = plt.subplot(n_models * 2, n_images, (i-1) * (n_images * 2) + n_images + j + 1)
        plt.imshow(decoded_imgs[j])
        plt.title("Reconstructed")
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

plt.show()
```



✓ ToDo : Testes (15pt)

Faça testes com vetor latente de dimensões 2, 8, 16 e 64.

Os testes foram realizados no próprio código acima mudando para imprimir todos os modelos com o vetor latente de dimensões diferentes.

✓ ToDo : Responda (15pt)

Escreva suas conclusões sobre os testes executados

Quanto maior for o vetor latente, melhor será a reconstrução da imagem.

✓ Parte II: Detecção de anomalias (30pt)

Intro

Neste exemplo, você vai detectar anomalias em sinais de eletrocardiograma (ECG). Para tal, treine um autoencoder no dataset [ECG5000 dataset](#). Este dataset contém 5000 batimentos de ECG (<https://en.wikipedia.org/wiki/Electrocardiography>), cada um com 140 amostras (pontos) na curva. Cada instância da base de dados (um batimento) foi rotulado como zero (0) ou um (1). A classe zero corresponde a um batimento anormal e a classe um a um batimento de classe normal. Queremos identificar os anormais.

Para detectar anomalias usando um autoencoder você deve treinar um autoencoder apenas em batimentos normais. Ele vai aprender a reconstruir os batimentos saudáveis. A hipótese é que os batimentos anormais vão divergir no padrão, quando compararmos a entrada com a re-construção.

✓ Carrega base de ECG

Base de dados detalhada no site: timeseriesclassification.com.

```
# Download the dataset
dataframe = pd.read_csv('http://storage.googleapis.com/download.tensorflow.org/data/ecg.csv', header=None)
raw_data = dataframe.values
dataframe.head()
```

	0	1	2	3	4	5	6	7	
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376041	-3.474986	-2.181408	-1.818286	-
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.992258	-1
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187449	-3.151462	-1.742940	-1.490659	-
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.671131	-
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423	-

5 rows × 141 columns

```
# The last element contains the labels
labels = raw_data[:, -1]

# The other data points are the electrocardiogram data
data = raw_data[:, 0:-1]

train_data, test_data, train_labels, test_labels = train_test_split(
    data, labels, test_size=0.2, random_state=21
)
```

Normaliza entre $[0,1]$.

```
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

Vamos separar os batimentos normais (label 1) para treinar o Autoencoder.


```
train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

normal_train_data = train_data[train_labels]
normal_test_data = test_data[test_labels]

anomalous_train_data = train_data[~train_labels]
anomalous_test_data = test_data[~test_labels]
```


Plote um batimento normal.

```
plt.grid()
plt.plot(np.arange(140), normal_train_data[0])
plt.title("A Normal ECG")
plt.show()
```



Plote um batimento anômalo.

```
plt.grid()
plt.plot(np.arange(140), anomalous_train_data[0])
plt.title("An Anomalous ECG")
plt.show()
```



✓ ToDo : Construção de um modelo (30pt)

Construa um modelo. Primeiramente tente construir apenas com camadas densas. Depois, tente construir um modelo com camadas de convolução de uma dimensão (Lembre-se que um sinal de ECG é uma série temporal de uma dimensão). [Conv1D](#)

```
class AnomalyDetector(Model):
    def __init__(self):
        super(AnomalyDetector, self).__init__()
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(32, activation='relu'),
            layers.Dense(16, activation='relu'),
            layers.Dense(8, activation='relu'),
        ])

        self.decoder = tf.keras.Sequential([
            layers.Dense(16, activation='relu'),
            layers.Dense(32, activation='relu'),
            layers.Dense(140, activation='sigmoid'), # Output size should match input size
            layers.Reshape((140, 1))
        ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = AnomalyDetector()
```

```
autoencoder.compile(optimizer='adam', loss='mae')
```

Depois de treinar com os batimentos normais, avalie com os anormais.

```
history = autoencoder.fit(normal_train_data, normal_train_data,
    epochs=20,
    batch_size=512,
    validation_data=(test_data, test_data),
    shuffle=True)
```

```
Epoch 1/20
5/5 [=====] - 2s 46ms/step - loss: 0.0585 - val_loss: 0.0537
Epoch 2/20
5/5 [=====] - 0s 11ms/step - loss: 0.0568 - val_loss: 0.0526
Epoch 3/20
5/5 [=====] - 0s 10ms/step - loss: 0.0553 - val_loss: 0.0514
Epoch 4/20
```

```

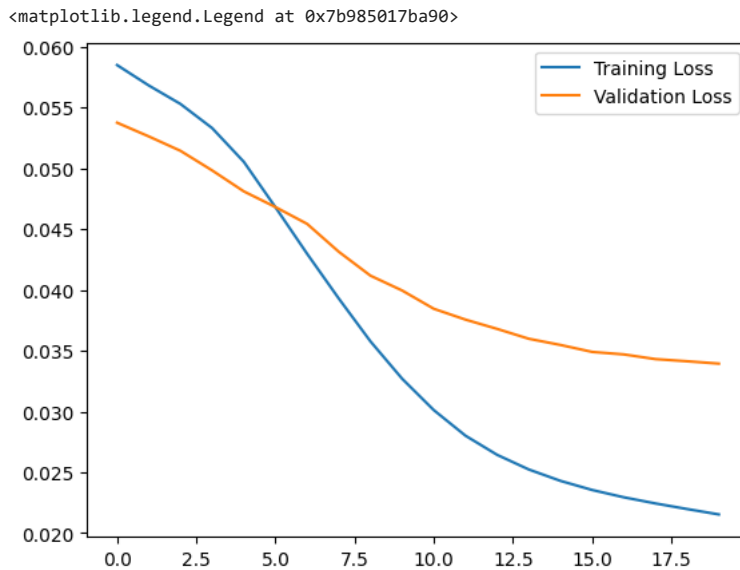
5/5 [=====] - 0s 14ms/step - loss: 0.0533 - val_loss: 0.0498
Epoch 5/20
5/5 [=====] - 0s 10ms/step - loss: 0.0505 - val_loss: 0.0481
Epoch 6/20
5/5 [=====] - 0s 10ms/step - loss: 0.0468 - val_loss: 0.0468
Epoch 7/20
5/5 [=====] - 0s 10ms/step - loss: 0.0430 - val_loss: 0.0454
Epoch 8/20
5/5 [=====] - 0s 10ms/step - loss: 0.0393 - val_loss: 0.0431
Epoch 9/20
5/5 [=====] - 0s 10ms/step - loss: 0.0358 - val_loss: 0.0412
Epoch 10/20
5/5 [=====] - 0s 10ms/step - loss: 0.0327 - val_loss: 0.0399
Epoch 11/20
5/5 [=====] - 0s 10ms/step - loss: 0.0301 - val_loss: 0.0384
Epoch 12/20
5/5 [=====] - 0s 10ms/step - loss: 0.0280 - val_loss: 0.0376
Epoch 13/20
5/5 [=====] - 0s 10ms/step - loss: 0.0264 - val_loss: 0.0368
Epoch 14/20
5/5 [=====] - 0s 10ms/step - loss: 0.0252 - val_loss: 0.0360
Epoch 15/20
5/5 [=====] - 0s 10ms/step - loss: 0.0243 - val_loss: 0.0355
Epoch 16/20
5/5 [=====] - 0s 10ms/step - loss: 0.0235 - val_loss: 0.0349
Epoch 17/20
5/5 [=====] - 0s 14ms/step - loss: 0.0229 - val_loss: 0.0347
Epoch 18/20
5/5 [=====] - 0s 10ms/step - loss: 0.0224 - val_loss: 0.0343
Epoch 19/20
5/5 [=====] - 0s 10ms/step - loss: 0.0220 - val_loss: 0.0341
Epoch 20/20
5/5 [=====] - 0s 10ms/step - loss: 0.0215 - val_loss: 0.0339

```

```

plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()

```



Você vai considerar um batimento como anômalo se ele divergir mais que um desvio padrão das amostras normais. Primeiro, vamos plotar um batimento normal a partir da base de treino e sua reconstrução. Assim, poderemos calcular o erro de re-construção.

```

encoded_data = autoencoder.encoder(tf.reshape(normal_test_data, (-1, 140))).numpy()
decoded_data = autoencoder.decoder(encoded_data).numpy()

plt.plot(normal_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')


decoded_data = np.squeeze(decoded_data)

plt.fill_between(np.arange(140), decoded_data[0], normal_test_data[0], color='lightcoral')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()

```




Imagem esperada:



Vamos fazer o mesmo para um batimento anômalo.

```
encoded_data = autoencoder.encoder(anomalous_test_data).numpy()
decoded_data = autoencoder.decoder(encoded_data).numpy()

plt.plot(anomalous_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')
decoded_data = np.squeeze(decoded_data)
plt.fill_between(np.arange(140), decoded_data[0], anomalous_test_data[0], color='lightcoral')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```





Imagem esperada:



✓ Detectando as anomalias

Vamos detectar as anomalias se o erro de reconstrução for maior que um limiar. Aqui, vamos calcular o erro médio para os exemplos normais do treino e depois, classificar os anormais do teste, que tenha erro de reconstrução maior que um desvio padrão.


Plota erro de reconstrução de batimentos normais do treino

```
reconstructions = autoencoder.predict(normal_train_data)
reconstructions = tf.squeeze(reconstructions, axis=-1)

train_loss = tf.keras.losses.mae(reconstructions, normal_train_data)

plt.hist(train_loss[None,:], bins=50)
plt.xlabel("Train loss")
plt.ylabel("No of examples")
plt.show()
```

74/74 [=====] - 0s 2ms/step



Escolha do limiar.


```
threshold = np.mean(train_loss) + np.std(train_loss)
print("Threshold: ", threshold)
```

Threshold: 0.03356406

```
reconstructions = autoencoder.predict(anomalous_test_data)
reconstructions = tf.squeeze(reconstructions, axis=-1)
test_loss = tf.keras.losses.mae(reconstructions, anomalous_test_data)

plt.hist(test_loss[None, :], bins=50)
plt.xlabel("Test loss")
plt.ylabel("No of examples")
plt.show()
```

14/14 [=====] - 0s 2ms/step



Classificação.

```
def predict(model, data, threshold):
    reconstructions = model(data)
    loss = tf.keras.losses.mae(reconstructions, data)
    return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
```

Calcule a acurácia para os dois modelos (com camadas densas e convolucionais)

```
print("Dimensões de reconstructions:", reconstructions.shape)
print("Dimensões de data:", data.shape)
data_subset = data[:440, :]
```

```
preds = predict(autoencoder, test_data, threshold)
print_stats(preds, test_labels)
```

```
Dimensões de reconstructions: (440, 140)
Dimensões de data: (4998, 140)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-29-c6bda96bca7b> in <cell line: 6>()
```

```
4
5 #preds = predict(autoencoder, test_data, threshold)
----> 6 print_stats(preds, test_labels)
```

```
NameError: name 'preds' is not defined
```

PESQUISAR NO STACK OVERFLOW

✓ Parte III: Redes Generativas Adversariais (40pt)

Leia o tutorial sobre a pix2pix em [Tensorflow Tutorials](#). O pix2pix foi apresentado em [Image-to-image translation with conditional adversarial networks by Isola et al. \(2017\)](#) e se trata de uma rede generativa adversarial condicional para geração de fachadas de prédios condicionada a uma máscara representando a arquitetura. baixe o notebook do tutorial, estude e treine a GAN. Após o treinamento, construa você mesmo 3 máscaras (usando algum software de desenho) e faça uma inferência com a rede. Anexe no notebook a máscara e sua respectiva saída.

✓ ToDo : Fachadas de prédios (40pt)

```
# ToDo : Criar 3 máscaras e gerar 3 saídas com a pix2pix para o problema de fachadas de prédios.
```

```
#Importar TensorFlow e outras bibliotecas
import tensorflow as tf
```

```
import os
import pathlib
import time
import datetime

from matplotlib import pyplot as plt
from IPython import display
```

Carregar o conjunto de dados

```
dataset_name = "facades"

_URL = f'http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/{dataset_name}.tar.gz'

path_to_zip = tf.keras.utils.get_file(
    fname=f"{dataset_name}.tar.gz",
    origin=_URL,
    extract=True)

path_to_zip = pathlib.Path(path_to_zip)


PATH = path_to_zip.parent/dataset_name

list(PATH.parent.iterdir())

sample_image = tf.io.read_file(str(PATH / 'train/1.jpg'))
sample_image = tf.io.decode_jpeg(sample_image)
print(sample_image.shape)

plt.figure()
plt.imshow(sample_image)
```

Downloading data from <http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades.tar>
 30168306/30168306 [=====] - 10s 0us/step
 (256, 512, 3)
 <matplotlib.image.AxesImage at 0x7b985e537ca0>



Defina uma função que carrega arquivos de imagem e gera dois tensores de imagem:

```
def load(image_file):
    # Read and decode an image file to a uint8 tensor
    image = tf.io.read_file(image_file)
    image = tf.io.decode_jpeg(image)

    # Split each image tensor into two tensors:
    # - one with a real building facade image
    # - one with an architecture label image
    w = tf.shape(image)[1]
    w = w // 2
    input_image = image[:, w:, :]
    real_image = image[:, :w, :]

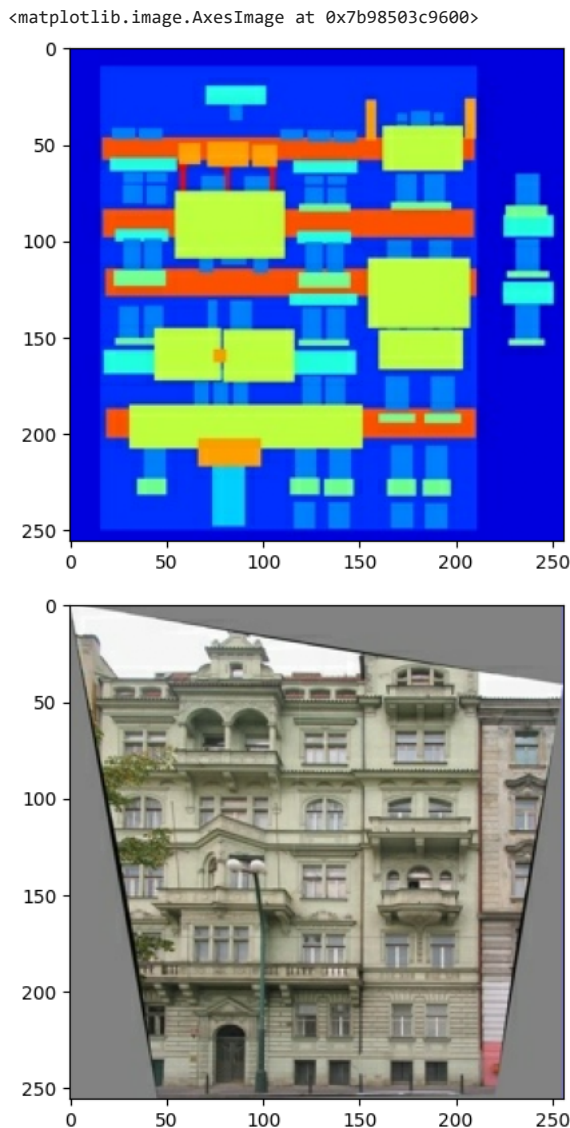
    # Convert both images to float32 tensors
    input_image = tf.cast(input_image, tf.float32)
    real_image = tf.cast(real_image, tf.float32)

    return input_image, real_image
```

plotando imagem

```
inp, re = load(str(PATH / 'train/100.jpg'))
# Casting to int for matplotlib to display the images
```

```
plt.figure()
plt.imshow(inp / 255.0)
plt.figure()
plt.imshow(re / 255.0)
```



Clique duas vezes (ou pressione "Enter") para editar

```
# The facade training set consist of 400 images
BUFFER_SIZE = 400
# The batch size of 1 produced better results for the U-Net in the original pix2pix experiment
BATCH_SIZE = 1
# Each image is 256x256 in size
IMG_WIDTH = 256
IMG_HEIGHT = 256
```

Clique duas vezes (ou pressione "Enter") para editar

```
def resize(input_image, real_image, height, width):
    input_image = tf.image.resize(input_image, [height, width],
                                   method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    real_image = tf.image.resize(real_image, [height, width],
                                   method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)

    return input_image, real_image
```

Clique duas vezes (ou pressione "Enter") para editar

```
def random_crop(input_image, real_image):
    stacked_image = tf.stack([input_image, real_image], axis=0)
    cropped_image = tf.image.random_crop(
        stacked_image, size=[2, IMG_HEIGHT, IMG_WIDTH, 3])

    return cropped_image[0], cropped_image[1]
```

Clique duas vezes (ou pressione "Enter") para editar

```
# Normalizing the images to [-1, 1]
def normalize(input_image, real_image):
    input_image = (input_image / 127.5) - 1
    real_image = (real_image / 127.5) - 1

    return input_image, real_image
```

Clique duas vezes (ou pressione "Enter") para editar

```
@tf.function()
def random_jitter(input_image, real_image):
    # Resizing to 286x286
    input_image, real_image = resize(input_image, real_image, 286, 286)


    # Random cropping back to 256x256
    input_image, real_image = random_crop(input_image, real_image)

    if tf.random.uniform(()) > 0.5:
        # Random mirroring
        input_image = tf.image.flip_left_right(input_image)
        real_image = tf.image.flip_left_right(real_image)

    return input_image, real_image
```

Clique duas vezes (ou pressione "Enter") para editar

```
plt.figure(figsize=(6, 6))
for i in range(4):
    rj_inp, rj_re = random_jitter(inp, re)
    plt.subplot(2, 2, i + 1)
    plt.imshow(rj_inp / 255.0)
    plt.axis('off')
plt.show()
```



Clique duas vezes (ou pressione "Enter") para editar

```
def load_image_train(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = random_jitter(input_image, real_image)
    input_image, real_image = normalize(input_image, real_image)

    return input_image, real_image
```

Clique duas vezes (ou pressione "Enter") para editar

```
def load_image_test(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = resize(input_image, real_image,
                                    IMG_HEIGHT, IMG_WIDTH)
    input_image, real_image = normalize(input_image, real_image)

    return input_image, real_image
```

Clique duas vezes (ou pressione "Enter") para editar

```
train_dataset = tf.data.Dataset.list_files(str(PATH / 'train/*.jpg'))
train_dataset = train_dataset.map(load_image_train,
                                  num_parallel_calls=tf.data.AUTOTUNE)
train_dataset = train_dataset.shuffle(BUFFER_SIZE)
train_dataset = train_dataset.batch(BATCH_SIZE)
```

Clique duas vezes (ou pressione "Enter") para editar

```
try:
    test_dataset = tf.data.Dataset.list_files(str(PATH / 'test/*.jpg'))
except tf.errors.InvalidArgumentError:
    test_dataset = tf.data.Dataset.list_files(str(PATH / 'val/*.jpg'))
test_dataset = test_dataset.map(load_image_test)
test_dataset = test_dataset.batch(BATCH_SIZE)
```

```
OUTPUT_CHANNELS = 3
```

```
def downsample(filters, size, apply_batchnorm=True):
    initializer = tf.random_normal_initializer(0., 0.02)

    result = tf.keras.Sequential()
    result.add(
        tf.keras.layers.Conv2D(filters, size, strides=2, padding='same',
                                kernel_initializer=initializer, use_bias=False))

    if apply_batchnorm:
        result.add(tf.keras.layers.BatchNormalization())

    result.add(tf.keras.layers.LeakyReLU())

    return result
```

```
down_model = downsample(3, 4)
down_result = down_model(tf.expand_dims(inp, 0))
print (down_result.shape)

(1, 128, 128, 3)
```

```
def upsample(filters, size, apply_dropout=False):
    initializer = tf.random_normal_initializer(0., 0.02)

    result = tf.keras.Sequential()
    result.add(
        tf.keras.layers.Conv2DTranspose(filters, size, strides=2,
                                         padding='same',
                                         kernel_initializer=initializer,
                                         use_bias=False))

    result.add(tf.keras.layers.BatchNormalization())

    if apply_dropout:
        result.add(tf.keras.layers.Dropout(0.5))

    result.add(tf.keras.layers.ReLU())

    return result
```

```
up_model = upsample(3, 4)
up_result = up_model(down_result)
print (up_result.shape)
```

```
(1, 256, 256, 3)
```

```
def Generator():
    inputs = tf.keras.layers.Input(shape=[256, 256, 3])

    down_stack = [
        downsample(64, 4, apply_batchnorm=False), # (batch_size, 128, 128, 64)
        downsample(128, 4), # (batch_size, 64, 64, 128)
        downsample(256, 4), # (batch_size, 32, 32, 256)
        downsample(512, 4), # (batch_size, 16, 16, 512)
        downsample(512, 4), # (batch_size, 8, 8, 512)
        downsample(512, 4), # (batch_size, 4, 4, 512)
        downsample(512, 4), # (batch_size, 2, 2, 512)
        downsample(512, 4), # (batch_size, 1, 1, 512)
    ]

    up_stack = [
        upsample(512, 4, apply_dropout=True), # (batch_size, 2, 2, 1024)
        upsample(512, 4, apply_dropout=True), # (batch_size, 4, 4, 1024)
        upsample(512, 4, apply_dropout=True), # (batch_size, 8, 8, 1024)
        upsample(512, 4), # (batch_size, 16, 16, 1024)
        upsample(256, 4), # (batch_size, 32, 32, 512)
        upsample(128, 4), # (batch_size, 64, 64, 256)
        upsample(64, 4), # (batch_size, 128, 128, 128)
    ]

    initializer = tf.random_normal_initializer(0., 0.02)
    last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNELS, 4,
                                           strides=2,
                                           padding='same',
                                           kernel_initializer=initializer,
                                           activation='tanh') # (batch_size, 256, 256, 3)

    x = inputs

    # Downsampling through the model
    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

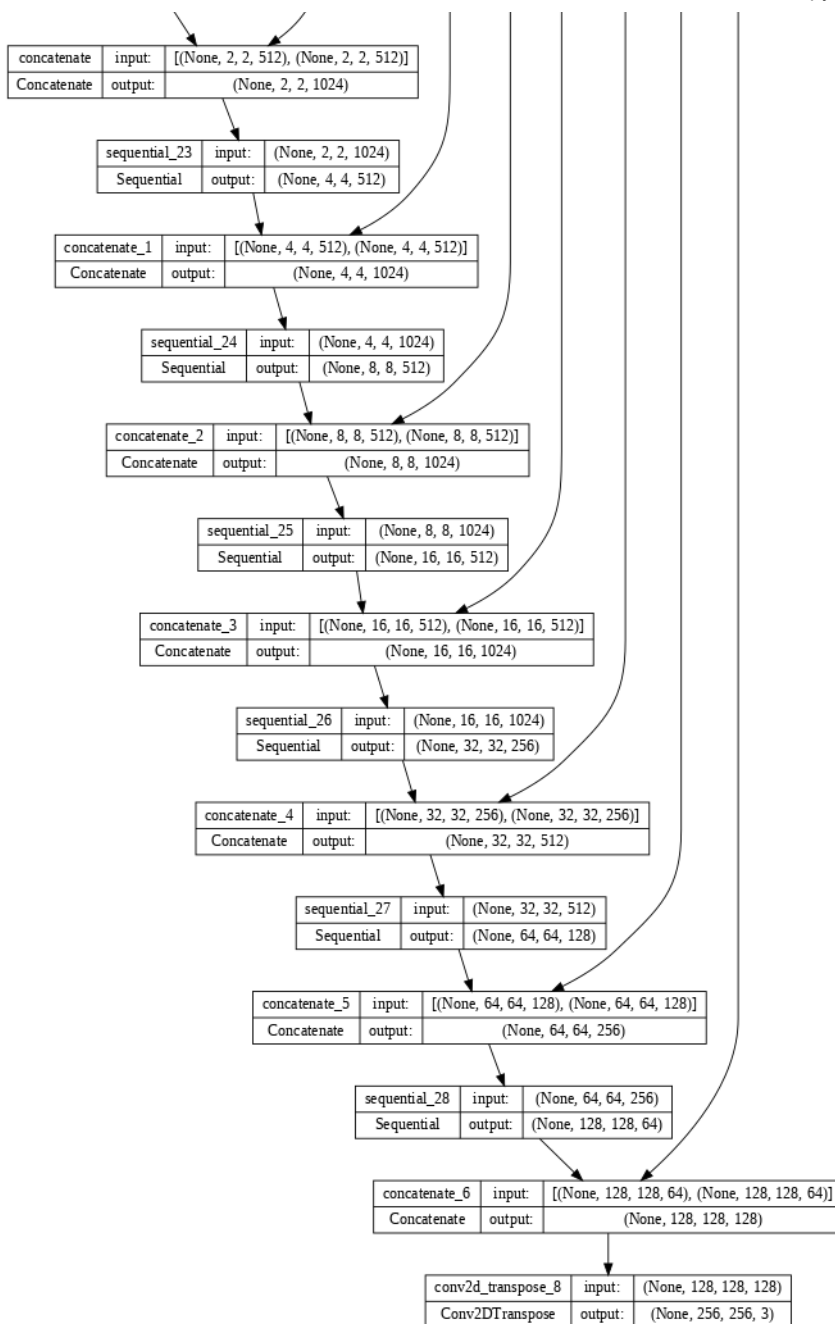
    skips = reversed(skips[:-1])

    # Upsampling and establishing the skip connections
    for up, skip in zip(up_stack, skips):
        x = up(x)
        x = tf.keras.layers.Concatenate()([x, skip])

    x = last(x)


    return tf.keras.Model(inputs=inputs, outputs=x)
```

```
generator = Generator()
tf.keras.utils.plot_model(generator, show_shapes=True, dpi=64)
```

```
gen_output = generator(inp[tf.newaxis, ...], training=False)
plt.imshow(gen_output[0, ...])
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB d
<matplotlib.image.AxesImage at 0x7b9832e89390>



LAMBDA = 100

```
loss_object = tf.keras.losses.BinaryCrossentropy(from_logits=True)
```

```
def generator_loss(disc_generated_output, gen_output, target):
    gan_loss = loss_object(tf.ones_like(disc_generated_output), disc_generated_output)

    # Mean absolute error
    l1_loss = tf.reduce_mean(tf.abs(target - gen_output))

    total_gen_loss = gan_loss + (LAMBDA * l1_loss)

    return total_gen_loss, gan_loss, l1_loss
```

```
def Discriminator():
    initializer = tf.random_normal_initializer(0., 0.02)

    inp = tf.keras.layers.Input(shape=[256, 256, 3], name='input_image')
    tar = tf.keras.layers.Input(shape=[256, 256, 3], name='target_image')

    x = tf.keras.layers.concatenate([inp, tar]) # (batch_size, 256, 256, channels*2)

    down1 = downsample(64, 4, False)(x) # (batch_size, 128, 128, 64)
    down2 = downsample(128, 4)(down1) # (batch_size, 64, 64, 128)
    down3 = downsample(256, 4)(down2) # (batch_size, 32, 32, 256)

    zero_pad1 = tf.keras.layers.ZeroPadding2D()(down3) # (batch_size, 34, 34, 256)
    conv = tf.keras.layers.Conv2D(512, 4, strides=1,
                                   kernel_initializer=initializer,
                                   use_bias=False)(zero_pad1) # (batch_size, 31, 31, 512)

    batchnorm1 = tf.keras.layers.BatchNormalization()(conv)

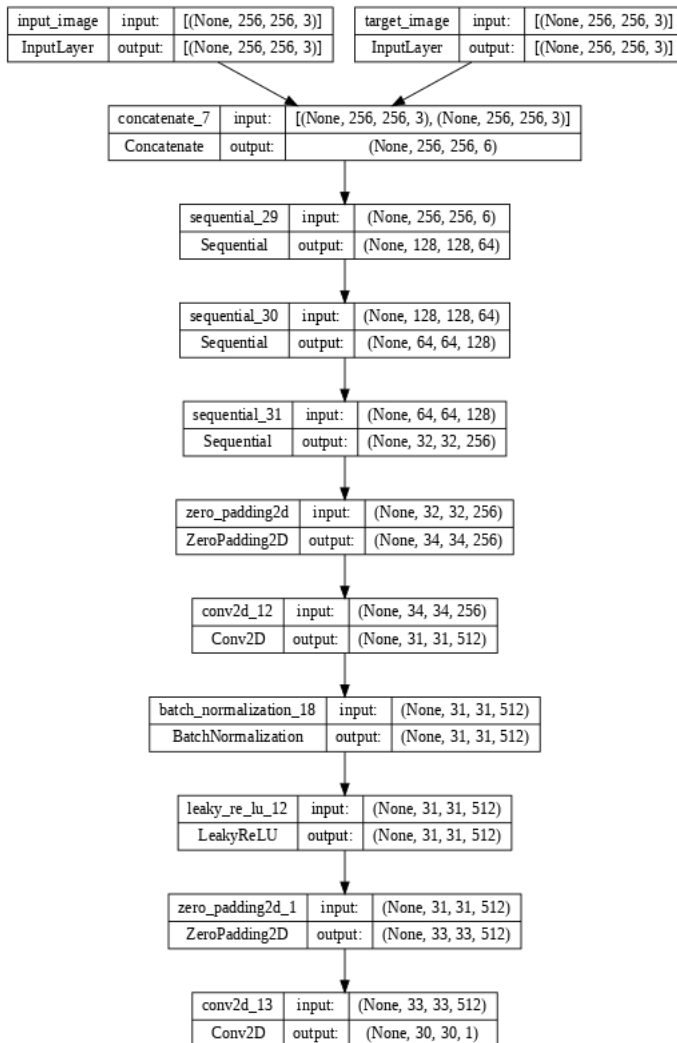
    leaky_relu = tf.keras.layers.LeakyReLU()(batchnorm1)

    zero_pad2 = tf.keras.layers.ZeroPadding2D()(leaky_relu) # (batch_size, 33, 33, 512)

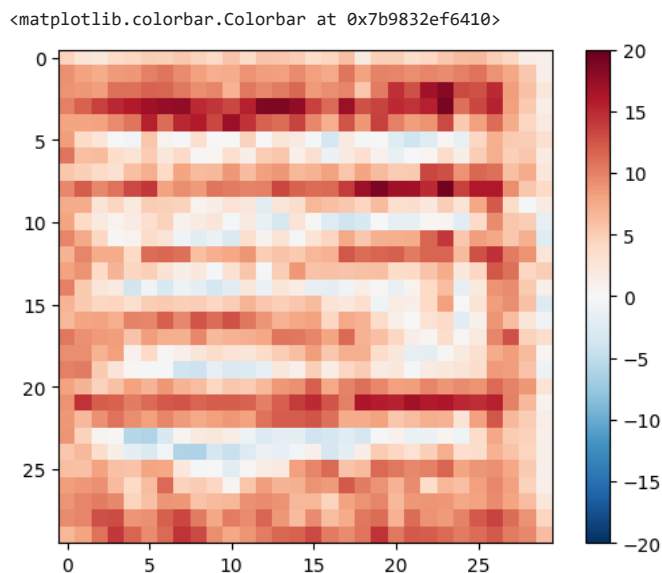
    last = tf.keras.layers.Conv2D(1, 4, strides=1,
                                   kernel_initializer=initializer)(zero_pad2) # (batch_size, 30, 30, 1)

    return tf.keras.Model(inputs=[inp, tar], outputs=last)
```

```
discriminator = Discriminator()
tf.keras.utils.plot_model(discriminator, show_shapes=True, dpi=64)
```



```
disc_out = discriminator([inp[tf.newaxis, ...], gen_output], training=False)
plt.imshow(disc_out[0, ..., -1], vmin=-20, vmax=20, cmap='RdBu_r')
plt.colorbar()
```



```
def discriminator_loss(disc_real_output, disc_generated_output):
    real_loss = loss_object(tf.ones_like(disc_real_output), disc_real_output)

    generated_loss = loss_object(tf.zeros_like(disc_generated_output), disc_generated_output)

    total_disc_loss = real_loss + generated_loss

    return total_disc_loss
```

```
generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
```

```
checkpoint_dir = './training_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                  discriminator_optimizer=discriminator_optimizer,
                                  generator=generator,
                                  discriminator=discriminator)
```

```
def generate_images(model, test_input, tar):
    prediction = model(test_input, training=True)
    plt.figure(figsize=(15, 15))

    display_list = [test_input[0], tar[0], prediction[0]]
    title = ['Input Image', 'Ground Truth', 'Predicted Image']

    for i in range(3):
        plt.subplot(1, 3, i+1)
        plt.title(title[i])
        # Getting the pixel values in the [0, 1] range to plot.
        plt.imshow(display_list[i] * 0.5 + 0.5)
        plt.axis('off')
    plt.show()
```

```
for example_input, example_target in test_dataset.take(1):
    generate_images(generator, example_input, example_target)
```